

CS341 HW2

Ryan Scherbarth

October 2024

1. **What is the assembly instruction to move a double word? How many bits does it move?**

movl is the instruction to move a double word. A word is 16 bits in x86-64, so it would move a total of 32 bits for a double word.

2. **What is the resulting memory address for the following operands assuming the indicated values in the registers?**

- (a) **-1(%rsp, %rcx, 4)**

`%rsp` = 0x09

`%rcx` = 0x05

Addr = `%rsp` + `%rcx` * 4 - 1

Addr = 0x21

- (b) **4(%rsp, %rcx, 4)**

`%rsp` = 0x08

`%rcx` = 0x0A

Addr = `%rsp` + `%rcx` * 4 + 4

Addr = 0x34

- (c) **256(, %rcx, 2)**

`%rcx` = 0x0A

Addr = `%rcx` * 2 + 256

Addr = 0x114

- (d) **0x33(, %rcx, 3)**

`%rcx` = 0x0B

Addr = `%rcx` * 3 + 0x33

Addr = 0x54

3. **What is the value in hexadecimal of the high-order 32 bits of register `%rax` after executing each of the following instructions? Provide an answer for each specific instruction.**

- (a) **movl \$0x11, %eax**

movl would take the value 0x11 and move it to the `%eax` register, which represents the lower 32 bits of the larger 64 bit `%rax` register. After the movl command is executed, moving the value to the lower order, the value of the higher order bits would simply be 0x00000000.

- (b) **movl \$0x0B, %eax**

Similar to part a, we will also end up with 0x00000000. Again, we are manipulating values with the `%eax` register. Whenever we use this register in the x86-86 bit OS it will automatically zero out the higher level bits.

4. **Provide the appropriate suffix and explain what is the instruction doing:**

`mov__ (%rax), %dx`

The complete instruction is: `movw (%rax), %dx`

`movw` introduces the `w` suffix, standing for word. Here, this means we are moving the size of a word, defined as 2 bytes (16 bits). `%rax` is the memory address, and `%dx` is the 16 bit register, which is the lower quarter (16 bits) of the `%rdx` register (64 bits).

5. **Consider these two declarations in C: `int *x` and `char *y`. Suppose that `x` is in register `%rdi` and `y` is in register `%rsi`, what is the assembly code to copy the value at the address stored in `%rdi` to the address stored in `%rsi`.**

To copy `x` from `%rdi` to `%rsi`, we have to start by de-referencing the pointer in `%rdi` (which will hold `int *`, a 4 byte value) and then copy it to the location `%rsi` is pointing to (which is holding a `char *`).

```
movl (%rdi), %eax
movl %eax, (%rsi)
```

Will properly copy the `int` value from the memory location that `%rdi` is pointing to into the memory location that `%rsi` is pointing to.

6. **What is the value in register `%rdx` after executing the following Assembly code**
`leaq -1(%rsp, %rcx, 4), %rdx`
assuming `%rsp` and `%rcx` have values `0x09` and `0x05`, respectively

`leaq`, load effective address, calculates the memory address but does not access that memory, so we are calculating the value and storing it in `%rdx`.

`%rsp` = `0x09`

`%rcx` = `0x05 * 4 = 0x14`

`%rdx` = `0x09 + 0x14 - 1 = 0x1C`

7. **What is the value (in decimal) in register `%rsp` after executing the following instructions in Assembly code?**

```
movl $1, %esp
addq $2, %rsp
incq %rsp
salq $1, %rsp
```

`movl $1, %esp` moves value 1 into `%esp` (the lower 32 bits of `%rsp`). The upper 32 bits of `%rsp` will be set to 0 whenever we write to `%esp`.

`addq $2, %rsp` is going to add 2 to the 64 bit register `%rsp`, so the result will now be `%rsp = 1+2 = 3`.

`incq %rsp` increments the register `%rsp` by 1, so the new value will be `%rsp = 3 + 1 = 4`.

`salq $1 %rsp` will perform a left shift on `%rsp` by a single bit, which is equivalent to multiplying the decimal value by 2. The result will now be `%rsp = 4 * 2 = 8`.

Therefore, the final value will be `%rsp = 8`.

8. What is the value (in decimal) in register `%rax` after executing the following instructions in assembly code?

```
movl $4, %eax
imulq $8, %rax
sarq $2, %rax
andq $4, %rax
```

`movl $4, %eax` moves the value 4 into the `%eax` register, which is the lower 32 bits of `%rax`. The first 32 bits of `%rax` will be set to 0.

`imulq $8, %rax` multiplies the value stored in `%rax` by 8, so the new value of `%rax` will now be 32.

`sarq $2, %rax` performs a right shift by 2 bits on the value stored in `%rax`. The right shift operation will preserve the sign bit, so after shifting the current value, 32, by 2 bits will result in the new value of `%rax` being 8.

`andq $4, %rax` performs a bitwise AND between `%rax` and 4. Converting both to binary, we have 1000 AND 0100. This will result in 0000 since there are no bits the same between the two inputs to the AND operation. The final value stored in `%rax` will be 0.

9. What is the value of the SF condition code after executing the following instructions in assembly code?

```
movb $0, %b1
cmpb $0x01, %b1
```

`movb $0, %b1` moves the value 0 into the 8 bit register `%b1`.

`cmpb $0x01, %b1` compares the 0x01 with the value in `%b1`, both of which are equal. The value of the SF flag is determined by comparing 0x01 with 0, and since 0 - 1 is negative, the SF flag gets set to 1 here.

10. What is the value of the ZF condition code after executing the following instructions in assembly code?

```
movb $1, %b1
cmpb $0x01, %b1
```

`movb $0, %b1` moves the value 0 into the 8 bit register `%b1`.

`cmpb $0x01, %b1` compares the 0x01 with the value in `%b1`, both of which are equal. Here, the ZF condition code will be set to 1 if the result of a subtraction is 0. When we do `cmpb` we are essentially just subtracting from 1 here, so the result of the ZF flag after this instruction will be set to 1.

11. Problem 3.63, p. 314-316

This problem will give you a chance to reverse engineer a switch statement from disassembled machine code. In the following procedure, the body of the switch statement has been omitted.

The jump table resides in a different area of memory. We can see from the indirect jump on line 5 that the jump table begins at address 0x4006f8. Using the GDB debugger, we can examine the six 8-byte words of memory comprising the jump table with the command `x/6gx 0x4006f8`. GDB prints the following:

```
1 (gdb) x/6gx 0x4006f8
2 0x4006f8:      0x00000000004005a1      0x00000000004005c3
3 0x400708:      0x00000000004005a1      0x00000000004005aa
4 0x400718:      0x00000000004005b2      0x00000000004005bf
```

Code:

```
1 long switch_prob(long x, long n){
2     long result = x;
3     switch(n) {
4         /* Fill in code here */
5     }
6     return result;
7 }
```

Disassembled Code:

```
1 0000000000400590 <switch_prob>
2 400590: 48 83 ee 3c          sub     $0x3c, %rsi
3 400594: 48 83 fe 05          cmp     $0x5, %rsi
4 400598: 77 29              ja      4005c3 <switch_prob+0x33>
5 40059a: ff 24 f5 f8 06 40 00 jmpq    *0x4006f8(,%rsi,8)
6 4005a1: 48 8d 04 fd 00 00 00 lea     0x0(,%rdi,8),%rax
7 4005a8: 00
8 4005a9: c3                retq
9 4005aa: 48 89 f8          mov     %rdi, %rax
10 4005ad: 48 c1 f8 03       sar     $0x3, %rax
11 4005b1: c3                retq
12 4005b2: 48 89 f8          mov     %rdi, %rax
13 4005b5: 48 c1 e0 04       shl     $0x4, %rax
14 4005b9: 48 29 f8          sub     %rdi, %rax
15 4005bc: 48 89 c7          mov     %rax, %rdi
16 4005bf: 48 0f af ff       imul    %rdi, %rdi
17 4005c3: 48 8d 47 4b       lea     0x4b(%rdi),%rax
18 4005c7: c3                retq
```

The default case is covered in lines 3-4

Here, we get another hint on line 3, where we compare n to 5, and then jump to a new location. This tells us there is a total of $[0..5] = 6$ cases including our base case.

case 0 and 2 correspond to lines 7-9

`lea` multiplies x by 8 and returns, so this case ends up just being $\text{result} = x * 8$.

case 1 corresponds to lines 6-8

Here, we move x to `%rax` and perform a right bit-wise shift of 3, which is equivalent to dividing by 8.

case 3 corresponds to lines 12-16

Here, we bitwise-shift x left by 4 (multiplying by 16), then subtract x , then we multiply this result by x . so we get $\text{result} = (x*16 - x) * x$

case 4 is given by lines 17-18

Here we load $x + 75$ into result, then we return $\text{result} = x + 75$ when $\text{case } n == 4$ is satisfied.

case 5 corresponds to lines 19-21

Here, we simply return $\text{result} = 12$

We can now reconstruct the switch case based off of our findings.

We can use the jump table on page 315 to determine the memory locations of the next entries in the jump table by looking at the output from the GDB command. The next entries (memory addresses) will be

2. 0x400700

4. 0x400710

8. 0x400718

```
switch(n) {
    case 0:
    case 2:
        result = 8x;
        break;
    case 1:
        result = x/8;
        break;
    case 3:
        result = (16x - x) * x;
        break;
    case 4:
        result = x + 75;
        break;
    case 5:
        result = 12;
        break;
    default:
        result = x + 75;
}
```

12. **What is the relationship between T_{\max} and U_{\max} ? Explain** an n bit signed int has a maximum value $T_{\max} = 2^{n-1} - 1$, since the other $n - 1$ bits are used for the magnitude, and one bit is saved for the sign.

an n bit unsigned integer, however, is given by $U_{\max} = 2^n - 1$, this is because we use all n bits to represent the magnitude in this case.

So we can see the relationship between the two numerically is $U_{\max} = 2 * T_{\max} + 1$

13. **Convert the value 0xF1AB to binary and apply the Binary two's complement (B2T) encoding to it, what is the resulting number?**

0xF = 1111

0x1 = 0001

$0xA = 1010$
 $0xB = 1011$
 $0xF1AB = 1111\ 0001\ 1010\ 1011$
 $0xF1AB_{inv} = 0000\ 1110\ 0101\ 0100$
 $0xF1AB_2's\ comp = 0000\ 1110\ 0101\ 0101$

14. **What is the value of x? (do it by hand, and show partial results at each step; do not run it in C)** `int x = (0xD2 & (1 << 7)) >> 7;`

$0xD2 = 1101\ 0010$
 $1101\ 0010 \& 1000\ 0000 = 1000\ 0000$
 $1000\ 0000 \gg 7 = 0000\ 0001$
 $x = 1$

15. **What is the value of x? (do it by hand, and show partial results at each step; do not run it in C)** `int x = (0xD2 & (1 << 3)) >> 3;`

$0xD2 = 1101\ 0010$
 $1101\ 0010 \& 0000\ 1000 = 0000\ 0000$
 $0000\ 0000 \gg 3 = 0000\ 0000$
 $x = 0$

16. Consider the 12-bit IEEE floating-point representation with 4 bits of exponent and 7 bits of fraction.

- (a) What is the bias? Show your intermediate work to compute the bias.

The bias is given by $B = 2^{(k-1)} - 1$.

$$\begin{aligned} B &= 2^{(4-1)} - 1 \\ B &= 2^{(4-1)} - 1 \\ &= 2^3 - 1 \\ &= 7 \end{aligned}$$

So we get a Bias of 7.

- (b) What is the value of 0111 1000 0000 in this 12-bit representation? Explain all the components of the bit pattern and which of the three cases it falls into.

In this 12-bit representation we have 1 sign bit, of the 10 remaining we know the bias to be 7, so the exponent is $12 - 1 - 7 = 4$.

sign: 1, exp: 4, mantissa: 7

The sign bit is 0, which means the number is positive.

the exponent bits are 1111, we can compute the exponent by subtracting the bias;

$$1111_2 = 15_{10}, \text{Exp} = 15 - 7 = 8$$

Since our exponent is not all 0's or 1's, we have a normalized representation, so the total value will be give by

$$0111 \ 1000 \ 0000 = 1.0 \times 2^8 = 256$$

- (c) What is the bit pattern for the decimal value 34.75 in this format?

The binary representation of the integer portion is $34_2 = 0010 \ 0010$.

The binary representation of the fraction will be 0.11 since $0.75 = \frac{1}{2^1} + \frac{1}{2^2}$

Now we have a direct binary representation of 0010 0010.11.

We next need to normalize our result. To normalize we will use the function $1.f \times 2^{\text{exp}}$
 $0010 \ 0010.11 \times 2^5$

The exponent will be given by the sum of the exponent and the bias; $5 + 127 = 132_{10} = 1000 \ 0100_2$.

Finally, we can combine all of the pieces to form our final number;

0 1000 0100 0000 1011 0000 0000 0000 000

17. Given the following encoding of a single-precision floating point number;

1 1000 0000 1100 0000 0000 0000 0000 000 ,

What are the following values? (Show all values in decimal; show M in both binary and decimal; show full equation for V)

- (a) $\text{exp} = 1000 \ 0000_2 = 128_{10}$
- (b) Bias = 127
- (c) $E = 128 - 127 = 1$
- (d) $M = 1 + 2^{-1} + 2^{-2} = 1.75$
- (e) $V = (-1)^S * M * 2^E = -3.5$

18. Given the following encoding of a single-precision floating point number;
1 1000 0001 1100 0000 0000 0000 0000 000 ,
 What are the following values? (Show all values in decimal; show M in both binary and decimal; show full equation for V)
- (a) $\text{exp} = 1000\ 0001_2 = 129_{10}$
 - (b) $\text{Bias} = 127$
 - (c) $E = 129 - 127 = 2$
 - (d) $M = 1 + 2^{-1} + 2^{-2} = 1.75$
 - (e) $V = -1^1(1.75 * 2^2) = -7$
19. Assuming Little Endian and that you have the addresses 0x404 to 0x407 available, how would the following data be stored in memory, if you start at address 00x404 0x5566AAA19?
- Our least significant bit will be stored at the lowest memory address, so to break up the data given;
- 0x404:** 0x19
0x405: 0xAA
0x406: 0x66
0x407: 0x55
- So the final result will be 0x19 0xAA 0x66 0x55
20. Assuming Little Endian and that you have the addresses 0x404 to 0x407 available, how would the following data be stored in memory, if you start at address 00x404 0x1CFF22?
- 0x404:** 0x22
0x405: 0xFF
0x406: 0x1C
0x407: 0x00

The final result will be 0x22 0xFF 0x1C