

# Lecture 3

## Preliminaries: Floating Point, Finite Precision, Significant digits

Owen L. Lewis

Department of Mathematics and Statistics  
University of New Mexico

August 27, 2024

# Recap: Absolute vs. Relative Error

- Here we used the absolute error:

$$\begin{aligned}\text{Error}(x_a) &= |x_t - x_a| \\ &= |\text{true value} - \text{approximate value}|\end{aligned}$$

- This doesn't tell the whole story. For example, if the values are large, like billions, then an Error of 100 is small. If the values are smaller, say around 10, then an Error of 100 is large. We need the relative error:

$$\begin{aligned}\text{Rel}(x_a) &= \left| \frac{x_t - x_a}{x_t} \right| \\ &= \left| \frac{\text{true value} - \text{approximate value}}{\text{true value}} \right|\end{aligned}$$

# Other bases

- Base for our number system is arbitrary!
- binary ( $\beta = 2$ ), octal ( $\beta = 8$ ), hexadecimal ( $\beta = 16$ ), etc
- In the  $\beta$ -system we have

$$(a_n \dots a_2 a_1 a_0 . b_1 b_2 b_3 b_4 \dots)_\beta = \sum_{k=0}^n a_k \beta^k + \sum_{k=0}^{\infty} b_k \beta^{-k}$$

# Changing Base

Base is arbitrary

We converted 325 to base 2.

Remember: Divide by 2 and record the remainders!

$(325)_{10}$  and  $(101000101)_2$  are the same number!

# Conversion from base-2 to base-10

Convert a base-2 number to base-10:

$$\begin{aligned}(1001.101)_2 \\&= 1 \times 2^3 + 0 \times 2^2 + 0 \times 2^1 + 1 \times 2^0 + 1 \times 2^{-1} + 0 \times 2^{-2} + 1 \times 2^{-3} \\&= 8 + 0 + 0 + 1 + 0.5 + 0 + 0.125 \\&= 9.62\end{aligned}$$

For other numbers, such as  $\frac{1}{5} = 0.2$ , an infinite length is needed.

$$0.2 \rightarrow .0011\,0011\,0011\,\dots$$

So 0.2 is stored just fine in base-10, but needs infinite number of digits in base-2

!!!

This is *roundoff* error in its basic form...

# Business Day

The New York Times

THURSDAY, NOVEMBER 24, 1994

Copyright © 1994 The New York Times

## BUSINESS Digest

THURSDAY, NOVEMBER 24, 1994

DOW	DOLLAR	OIL	BONDS
30	vs. Japanese	Midwest	30-Year
Industrials	Yen	Spot	Treasury
3,674.83	98.47 Yen	\$18.15	7.94%
-3.36	+0.24 Yen	+\$0.33	-0.08

### Companies

A flaw in the Pentium, the top chip made by Intel, can cause inaccurate calculations in certain rare cases, Intel said. The chip would not have to be recalled, but many scientists and engineers who depend on precise calculations are concerned. [Page D1.]

Gibson Greetings and Bankers Trust settled a lawsuit over derivatives. Gibson agreed to pay Bankers Trust nearly \$8.2 million, or about 36 percent of the \$22.7 million that Bankers claimed it was owed in the trades between the two. [D1.]

Sony, Charlie H. J. Heller and Leo Burnett are parting ways after 27 years. The Burnett agency created familiar advertising characters like Morris the monkey cat and Charlie the tuna. [D1.]

Pfizer will buy SmithKline Beecham's animal health business for \$1.65 billion in cash. The deal would make Pfizer the world's largest maker of animal drugs, topping Merck. [D4.]

Kmart is seeking ways to save as much as \$800 million in two years in addition to store closings and job cuts. [D4.]

A Federal judge from a New York brokerage account of a British businessman accused of illegal insider trading. [D4.]

The F.T.C. is seeking data on alliances with Caraway International, in a broadening of an investigation of relations between drug makers and managed health-care drug distributors. [D4.]

Metallgesellschaft of Germany plans to write down the paper value of its shares by half in an effort to rescue itself. [D4.]

I.B.M., AT&T, Apple Computer and others plan to create a universal communication language for computers. [D5.]

### Markets

After a roller-coaster session, stocks ended slightly lower. The Standard & Poor's 500 index fell just 0.16 point, easing investors' concerns that stocks were in for a long, hard descent. [D1.]

A stock slide is not likely to change Federal Reserve policy. Indeed, stocks' strength until now surprised Fed officials. [D1.]

Stock markets fell sharply across Europe as investors took their cue from Wall Street and sought refuge in bonds. [D1.]

## Flaw Undermines Accuracy of Pentium Chips

By JOHN MARKOFF

Special to The New York Times

SAN FRANCISCO, Nov. 23 — An elusive circuitry error is causing a chip used in millions of computers to generate inaccurate results in certain rare cases, heightening anxiety among many scientists and engineers who rely on their machines for precise calculations.

The flaw, an error in division, has been found in the Pentium, the current top microprocessor of the Intel Corporation, the world's largest chip maker. The chip, in several different configurations, is used in many com-

puters sold for home and business use, including those made by I.B.M., Compaq, Dell, Gateway 2000 and others.

The flaw appears in all Pentium chips now on the market, in certain types of division problems involving more than five significant digits, a mathematical turn that can include numbers before and after a decimal point.

Intel declined to say how many Pentium chips it made or sold, but Dataquest, a market research company in San Jose, Calif., estimated that in 1994 Intel would sell 5.5 million to 6 million Pentiums, roughly 10 percent of the number of personal

computers sold worldwide.

Intel said yesterday that it did not believe the chip needed to be recalled, asserting that the typical user would have but one chance in more than nine billion of encountering an inaccurate result as a consequence of the error, and thus there was no noticeable consequence to users of business or home computers. Indeed, the company said it was continuing to send computer makers Pentium chips built before the problem was detected.

William Kahan of the University of California at Berkeley, one of the nation's experts on computer mathematics, expressed skepticism about

Intel's contention that the error would only occur in extremely rare instances.

"These kinds of statistics have to cause some wonderment," he said. "They are based on assertions about the probability of events whose probability we don't know."


At the Jet Propulsion Laboratory in Pasadena, Calif., one satellite communications researcher who learned of the error this week, said six Pentium machines were used in his group and their use had been suspended for now.

The Pentium appeared as a con-

Continued on Page D5

## Four Sharp Corrections in the Dow

The year the stock market has tended to gain gradually for weeks or had nervously around a trading range with little direction, waiting for the moment when some bad news or unfavorable trend sets off a heading flight. Then, once prices have been beaten down for a few days, the buyers reappear and the bleeding stops. Four of these corrections are highlighted, along with the development that traders said pushed the market over the precipice.



## Gibson Suit On Trades Is Settled

### Bankers Trust Gets 30% of Debt Claimed

By MICHAEL QUINT

Gibson Greetings Inc. and the Bankers Trust Company said yesterday that they had settled Gibson's lawsuit accusing the bank of improperly trading the company to engage in risky financial trades.

Under the out-of-court settlement, Gibson will pay Bankers Trust nearly \$8.2 million, or about 36 percent of the \$22.7 million that Bankers Trust contended it was owed under the trades.

While neither side boasted of victory, Douglas Kidd, a spokesman for Bankers Trust, acknowledged that

## Flaw Undermines Accuracy of Pentium Chips

By JOHN MARKOFF/Special to The New York Times

*New York Times (1857-Current file)*; Nov 24, 1994; ProQuest Historical Newspapers The New York Times (1851 - 2002)  
pg. D1

# Flaw Undermines Accuracy of Pentium Chips

By JOHN MARKOFF

Special to The New York Times

SAN FRANCISCO, Nov. 23 — An elusive circuitry error is causing a chip used in millions of computers to generate inaccurate results in certain rare cases, heightening anxiety among many scientists and engineers who rely on their machines for precise calculations.

The flaw, an error in division, has been found in the Pentium, the current top microprocessor of the Intel Corporation, the world's largest chip maker. The chip, in several different configurations, is used in many com-

puters sold for home and business use, including those made by I.B.M., Compaq, Dell, Gateway 2000 and others.

The flaw appears in all Pentium chips now on the market, in certain types of division problems involving more than five significant digits, a mathematical term that can include numbers before and after a decimal point.

Intel declined to say how many Pentium chips it made or sold, but Dataquest, a market research company in San Jose, Calif., estimated that in 1994 Intel would sell 5.5 million to 6 million Pentiums, roughly 10 percent of the number of personal

computers sold worldwide.

Intel said yesterday that it did not believe the chip needed to be recalled, asserting that the typical user would have but one chance in more than nine billion of encountering an inaccurate result as a consequence of the error, and thus there was no noticeable consequence to users of business or home computers. Indeed, the company said it was continuing to send computer makers Pentium chips built before the problem was detected.

William Kahan of the University of California at Berkeley, one of the nation's experts on computer mathematics, expressed skepticism about

Intel's contentions that the error would only occur in extremely rare instances.

"These kinds of statistics have to cause some wonderment," he said. "They are based on assertions about the probability of events whose probability we don't know."

At the Jet Propulsion Laboratory in Pasadena, Calif., one satellite communications researcher who learned of the error this week, said six Pentium machines were used in his group and their use had been suspended for now.

"The Pentium appeared as a cost-

*Continued on Page D5*

## tium Chips

In some complex division problems, annoying errors.

corrected.

Some computer users said they believed that Intel had not acted quickly enough after discovering the error.

"Intel has known about this since the summer; why didn't they tell anyone?" said Andrew Schulman, the author of a series of technical books on PC's. "It's a hot issue, and I don't think they've handled this well."

The company said that after it discovered the problem this summer, it ran months of simulations of different applications, with the help of outside experts, to determine whether the problem was serious.

The Pentium error occurs in a portion of the chip known as the floating point unit, which is used for extremely precise computations. In rare cases, the error shows up in the result of a division operation.

Intel said the error occurred because of an omission in the translation of a formula into computer

### Close, but Not Close Enough

The owners of computers that use Intel's Pentium microprocessors have found that the chips sometimes do not perform division calculations accurately enough.

The problems arise when the chip has to round a number in a preliminary calculation to get the final result, a task that all processors normally perform. In these cases, however, the Pentium's figures are exact to only 5 digits, not 16, as are those of other computer processors. The Pentium's error, while small, can be 10 billion times as large as those of most chips.

Here is an example of the way the imprecise rounding changes the results of a calculation and the way the deviation from the expected result is calculated.

#### PROBLEM

$$4,195,835 - [(4,195,835 + 3,145,727) \times 3,145,727]$$

#### CORRECT CALCULATION

$$= 4,195,835 - [(1.3338204) \times 3,145,727] = 0$$

#### PENTIUM'S CALCULATION

$$= 4,195,835 - [(1.3337391) \times 3,145,727] = 256$$

#### DEVIATION

$$256 \div 4,195,835 = 6.1 \times 10^{-5}, \text{ or } 61/100,000$$

Source: Cleve Moler, the Mathworks Inc.



# Intel Timeline

from emery.com

June 1994 Intel engineers discover the division error. Managers decide the error will not impact many people. Keep the issue internal.

June 1994 Dr Nicely at Lynchburg College notices computation problems

Oct 19, 1994 After months of testing, Nicely confirms that other errors are not the cause. The problem is in the Intel Processor.

Oct 24, 1994 Nicely contacts Intel. Intel duplicates error.

Oct 30, 1994 After no action from Intel, Nicely sends an email

# Intel Timeline

from emery.com

FROM: Dr. Thomas R. Nicely  
Professor of Mathematics  
Lynchburg College  
1501 Lakeside Drive  
Lynchburg, Virginia 24501-3199

Phone: 804-522-8374  
Fax: 804-522-8499  
Internet: nicely@acavax.lynchburg.edu

TO: Whom it may concern

RE: Bug in the Pentium FPU

DATE: 30 October 1994

It appears that there is a bug in the floating point unit (numeric coprocessor) of many, and perhaps all, Pentium processors.

In short, the Pentium FPU is returning erroneous values for certain division operations. For example,

$0001/824633702441.0$

is calculated incorrectly (all digits beyond the eighth significant digit are in error). This can be verified in compiled code, an ordinary spreadsheet such as Quattro Pro or Excel, or even the Windows calculator (use the scientific mode), by computing

$00(824633702441.0)*(1/824633702441.0)$ ,

which should equal 1 exactly (within some extremely small rounding error; in general, coprocessor results should contain 19 significant decimal digits). However, the Pentiums tested return

$0000.999999996274709702$

.  
.  
.

# Intel Timeline

from emery.com

- Nov 1, 1994 Software company Phar Lap Software receives Nicely's email. Sends to colleagues at Microsoft, Borland, Watcom, etc. decide the error will not impact many people. Keep the issue internal.
- Nov 2, 1994 Email with description goes global.
- Nov 15, 1994 USC reverse-engineers the chip to expose the problem. Intel still denies a problem. Stock falls.
- Nov 22, 1994 CNN *Moneyline* interviews Intel. Says the problem is minor.
- Nov 23, 1994 The MathWorks develops a fix.
- Nov 24, 1994 New York Times story. Intel still sending out flawed chips. Will replace chips only if it caused a problem in an important application.
- Dec 12, 1994 IBM halts shipment of Pentium based PCs
- Dec 16, 1994 Intel stock falls again.
- Dec 19, 1994 More reports in the NYT: lawsuits, etc.
- Dec 20, 1994 Intel admits. Sets aside \$420 million to fix.

# Numerical "bugs"

## Obvious

Software has bugs

## Not-SO-Obvious

Numerical software has two unique bugs:

- 1 roundoff error
- 2 truncation error

# Numerical Errors

## Roundoff

*Roundoff* occurs when digits in a decimal point (0.3333...) are lost (0.3333) due to a limit on the memory available for storing one numerical value.

## Truncation

*Truncation* error occurs when discrete values are used to approximate a mathematical expression.

# Uncertainty: well- or ill-conditioned?

Errors in input data can cause *uncertain* results

- input data can be experimental or rounded. leads to a certain variation in the results
- *well-conditioned*: numerical results are insensitive to small variations in the input
- *ill-conditioned*: small variations lead to drastically different numerical calculations (a.k.a. poorly conditioned)

# Our Job

As numerical analysts, we need to

- ① solve a problem so that the calculation is not susceptible to large roundoff error
- ② solve a problem so that the approximation has a *tolerable* truncation error

How?

- incorporate roundoff-truncation knowledge into
  - the mathematical model
  - the method
  - the algorithm
  - the software design
- awareness → correct interpretation of results

# Floating Points

## Normalized Floating-Point Representation

Real numbers are stored as

$$x = \pm(0.d_1d_2d_3\dots d_m)_\beta \times \beta^e$$

- $d_1d_2d_3\dots d_m$  is the mantissa,  $e$  is the exponent
- $e$  is negative, positive or zero
- the general normalized form requires  $d_1 \neq 0$



# Floating Point

## Example

In base 10

- 1000.12345 can be written as

$$(0.100012345)_{10} \times 10^4$$

- 0.000812345 can be written as

$$(0.812345)_{10} \times 10^{-3}$$

- $1/3$  can be written as

$$(0.33333\dots)_{10} \times 10^0$$

# Floating Point

Floating point is not specific to any base.

## Example

325 can be written as

$$(0.325)_{10} \times 10^3$$

or ...

$$(0.1101000101)_2 \times 2^9.$$

# Notes

# Floating Point (with finite precision)

Suppose we have only 3 bits for a mantissa and a 2 bit exponent stored like

$.d_1$	$d_2$	$d_3$	$e_1$	$e_2$
$(1/2)$	$(1/4)$	$(1/8)$	$(\pm)$	$1, 0$


All possible combinations give:

$$000_2 = 0$$

$$\dots \times 2^{-1,0,1}$$


$$111_2 = 7$$

So we get  $0, \frac{1}{16}, \frac{2}{16}, \dots, \frac{7}{16}, 0, \frac{1}{4}, \frac{2}{4}, \dots, \frac{7}{4}$ , and  $0, \frac{1}{8}, \frac{2}{8}, \dots, \frac{7}{8}$ . On the real line:



# Notes


# Overflow, Underflow



- computations too close to zero may result in *underflow*
- computations too large may result in *overflow*
- overflow error is considered more severe
- underflow can just fall back to 0

# Normalizing

If we use the normalized form in our 5-bit case, we lose  $[(0.001)_2, (0.010)_2, (0.011)_2] \times 2^{-1,0,1}$ . So we cannot represent  $\frac{1}{16}$ ,  $\frac{1}{8}$ , and  $\frac{3}{16}$ .




- IEEE-754 is a widely used standard accepted by hardware/software makers
  - defines the floating point distribution for our computation
  - offer several rounding modes which effect accuracy
- Floating point arithmetic emerges in nearly every piece of code
  - even modest mathematical operations yield loss of significant bits
  - several pitfalls in common mathematical expressions



# IEEE Floating Point (IEEE-754)


- How much storage do we actually use in practice?
- 32-bit word lengths are typical
- Single-precision floating-point numbers use 32 bits
- Double-precision floating-point numbers use 64 bits
- Long Double-precision floating-point numbers use 80 bits
- Goal: use the 32-bits to best represent the normalized floating point number



# Notes

# IEEE Single Precision

$$x = \pm q \times 2^m$$



## Notes:

- 1-bit sign
- 8-bit exponent  $|m|$
- 23-bit mantissa  $q$
- The leading one in the mantissa  $q$  does not need to be represented:  $b_1 = 1$  is hidden bit
- IEEE 754: put  $x$  in 1.f normalized form
- Largest exponent = 127, Smallest exponent = -126  
Achieved without a sign bit for the exponent:  $0 < m + 127 = c < 255$
- Special cases:  $c = 0, 255$  for  $\pm 0, \pm \text{inf}, \text{nan} \dots$
- use 1-bit sign for negatives (1) and positives (0)

# IEEE Single Precision

$$x = \pm q \times 2^m$$


①  $(-52.125)_{10} = -(110100.001000000000)_2$

② Convert to 1.f form:  $x = \underbrace{-}_1 (1. \underbrace{101\ 000\ 010\ 000\ \dots\ 0}_{23})_2 \times 2^5$

③ Convert exponent  $5 = c - 127 \Rightarrow c = 132 \Rightarrow c = (\underbrace{10\ 000\ 100}_8)_2$

$$\underbrace{1}_1 \underbrace{10\ 000\ 100}_8 \underbrace{101\ 000\ 010\ 000\ \dots\ 0}_{23}$$

# IEEE Double Precision



- 1-bit sign
- 11-bit exponent
- 52-bit mantissa
- single-precision: about 6 decimal digits of precision
- double-precision: about 15 decimal digits of precision
- $m = c - 1023$

# Precision vs. Range

type	range	approx range
single	$\begin{array}{c} -3.40 \times 10^{38} \leq x \leq -1.18 \times 10^{-38} \\ 0 \\ 1.18 \times 10^{-38} \leq x \leq 3.40 \times 10^{38} \end{array}$	$2^{-126} \rightarrow 2^{128}$
double	$\begin{array}{c} -1.80 \times 10^{308} \leq x \leq -2.23 \times 10^{-308} \\ 0 \\ 2.23 \times 10^{-308} \leq x \leq 1.80 \times 10^{308} \end{array}$	$2^{-1022} \rightarrow 2^{1024}$

Listing 1: Matlab Tools

```
1 >> realmax
2 ans = 1.7977e+308
3
4 >> realmin
5 ans = 2.2251e-308
```

# Notes

Take  $x = 1.0$  and add  $1/2, 1/4, \dots, 2^{-i}$ :

Hidden bit													
↓	←	52 bits										→	
1	1	0	0	0	0	0	0	0	0	0	0	e	e
1	0	1	0	0	0	0	0	0	0	0	0	e	e
1	0	0	1	0	0	0	0	0	0	0	0	e	e
.....													
1	0	0	0	0	0	0	0	0	0	0	1	e	e
1	0	0	0	0	0	0	0	0	0	0	0	e	e

- Ooops!
- Define  $fl(x)$  to represent the floating point machine number for the real number  $x$
- $fl(1 + 2^{-52}) \neq 1$ , but  $fl(1 + 2^{-53}) = 1$



## Machine Epsilon

The machine epsilon  $\epsilon_m$  is the smallest number such that


$$fl(1 + \epsilon_m) \neq 1$$

- The double precision machine epsilon is about  $2^{-52}$ .
- The single precision machine epsilon is about  $2^{-23}$ .

```
1 >> eps
2 ans = 2.2204e-16
3
4 >> eps('single')
5 ans = 1.1921e-07
```

“Realmin” and “machine epsilon” are not the same thing!

# Floating Point Number Line



# Floating Point Errors

- Not all reals can be exactly represented as a machine floating point number. Then what?
- Round-off error
- IEEE options:
  - Round to next nearest FP (preferred), Round to 0, Round up, and Round down

Let  $x_+$  and  $x_-$  be the two floating point machine numbers closest to  $x$

- round to nearest:  $round(x) = x_-$  or  $x_+$ , whichever is closest
- round toward 0:  $round(x) = x_-$  or  $x_+$ , whichever is between 0 and  $x$
- round toward  $-\infty$  (down):  $round(x) = x_-$
- round toward  $+\infty$  (up):  $round(x) = x_+$

# Floating Point Errors

How big is this error? Suppose ( $x$  is closer to  $x_-$ )

$$x = (1.b_2b_3 \dots b_{24}b_{25}b_{26})_2 \times 2^m$$

$$x_- = (1.b_2b_3 \dots b_{24})_2 \times 2^m$$

$$x_+ = ((1.b_2b_3 \dots b_{24})_2 + 2^{-24}) \times 2^m$$

$$|x - x_-| \leq \frac{|x_+ - x_-|}{2} = 2^{m-25} \quad : \text{absolute error}$$

$$\left| \frac{x - x_-}{x} \right| \leq \frac{2^{m-25}}{1/2 \times 2^m} \leq 2^{-24} = \epsilon_m/2 \quad : \text{relative error}$$


# Notes

# Floating Point Arithmetic

- Problem: The set of representable machine numbers is FINITE.
- So not all math operations are well defined!
- Basic algebra breaks down in floating point arithmetic

## Example

$$a + (b + c) \neq (a + b) + c$$



# Floating Point Arithmetic

Rule 1.

$$fl(x) = x(1 + \epsilon), \quad \text{where } |\epsilon| < \epsilon_m$$

Rule 2.

For all operations  $\odot$  (one of  $+$ ,  $-$ ,  $*$ ,  $/$ )

$$fl(x \odot y) = (x \odot y)(1 + \epsilon_{\odot}), \quad \text{where } |\epsilon_{\odot}| < \epsilon_m$$

Rule 3.

For  $+$ ,  $*$  operations

$$fl(a \odot b) = fl(b \odot a)$$

There were many discussions on what conditions/ rules should be satisfied by floating point arithmetic. The IEEE standard is a set of standards adopted by many CPU manufacturers.

# Floating Point Arithmetic

Consider the sum of 3 numbers:  $y = a + b + c$ .

Done as  $fl(fl(a + b) + c)$

$$\begin{aligned}\eta &= fl(a + b) = (a + b)(1 + \epsilon_1) \\ y_1 &= fl(\eta + c) = (\eta + c)(1 + \epsilon_2) \\ &= [(a + b)(1 + \epsilon_1) + c](1 + \epsilon_2) \\ &= [(a + b + c) + (a + b)\epsilon_1](1 + \epsilon_2) \\ &= (a + b + c) \left[ 1 + \frac{a + b}{a + b + c} \epsilon_1 (1 + \epsilon_2) + \epsilon_2 \right]\end{aligned}$$

So disregarding the high order term  $\epsilon_1 \epsilon_2$

$$fl(fl(a + b) + c) = (a + b + c)(1 + \epsilon_3) \quad \text{with} \quad \epsilon_3 \approx \frac{a + b}{a + b + c} \epsilon_1 + \epsilon_2$$



# Floating Point Arithmetic

If we redid the computation as  $y_2 = fl(a + fl(b + c))$  we would find

$$fl(a + fl(b + c)) = (a + b + c)(1 + \epsilon_4) \quad \text{with} \quad \epsilon_4 \approx \frac{b + c}{a + b + c} \epsilon_1 + \epsilon_2$$

Main conclusion:

The first error is amplified by the factor  $(a + b)/y$  in the first case and  $(b + c)/y$  in the second case.

In order to sum  $n$  numbers more accurately, it is better to start with the small numbers first. [However, sorting before adding is usually not worth the cost!]

# Significant Digits

## Significant Digits

Digits in a number justified by the accuracy of the means of measurement

Example: Measurement from a ruler which goes from 0.0 to 9.0 (in cm):

- $x = .23 \times 10^1 \text{ cm}$
- two significant digits
- 2 is the most significant digit
- 3 is the least significant digit
  
- $x = .40 \times 10^1 \text{ cm}$
- still two significant digits