

Lecture 13

Symmetric Factorizations & Iterative Methods

Owen L. Lewis

Department of Mathematics and Statistics
University of New Mexico

Oct. 1, 2024

Goals for today. . .

- Review LDL^T factorization
- Show Cholesky factorization.
- Discussion of Matrix solve costs.
- Iterative methods

Reminder: Goal of Factorization?

- Different solves have different costs:
 - $Ax = b$ costs $\mathcal{O}(n^3)$ via Gaussian elimination.
 - $Ly = b$ or $Uy = b$ costs $\mathcal{O}(n^2)$ via forward/backward substitution.
 - $\rightarrow A = LU$ costs $\mathcal{O}(n^3)$ to compute L and U (basically just Gaussian Elimination), but each successive solve is just 2 triangular solves.
- What if we know that $A = A^T$, can we speed up computation of L and U ?

Use SYMMETRY ! YRTEMMYS esU

Suppose

$$A = A^T.$$

$$\Rightarrow LU = A = LDL^T.$$

LDL^T Factorization

If A is symmetric, we can factor A as

$$A = LDL^T,$$

where L is the same lower triangular matrix made during LU factorization (don't need to do any work to create L^T).

How do we make D ?

If you multiply out $A = LDL^T$ diagonal by diagonal and take advantage of all the zeros, you get

$$d_{11} = a_{11},$$

$$d_{22} = a_{22} - d_{11}\ell_{21}^2,$$

$$d_{33} = a_{33} - d_{11}\ell_{31}^2 - d_{22}\ell_{32}^2,$$

$$\vdots$$

LDL^T Factorization

Listing 1: LDL^T Factorization

```
1  given  $A$   
2  output  $L, D$   
3  
4  for  $j = 1 \dots n$   
5       $\ell_{jj} = 1$   
6  
7       $d_j = a_{jj} - \sum_{v=1}^{j-1} d_v \ell_{jv}^2$ . % COULD BE NEGATIVE  
8  
9      for  $i = j + 1 \dots n$   
10          $\ell_{ji} = 0$   
11          $\ell_{ij} = \left( a_{ij} - \sum_{v=1}^{j-1} \ell_{iv} d_v \ell_{jv} \right) / d_j$   
12     end  
13 end
```

- Special form of the LU factorization (for symmetric A).

$\tilde{L}\tilde{L}^T$: Cholesky Factorization

- A must be symmetric and positive definite (SPD)
- A is Positive Definite (PD) if for all $x \neq 0$ the following holds

$$x^T A x > 0$$

- Positive definite gives us an all positive D in $A = LDL^T$
- Define $D^{1/2}$ via $\sqrt{d_{ii}}$.
- $A = LDL^T = LD^{1/2}D^{1/2}L^T$.
- Define $\tilde{L} = LD^{1/2}$.
- $A = \tilde{L}\tilde{L}^T$, i.e. $\tilde{L} = \tilde{U}^T$ (note, L had ones on the diagonal, \tilde{L} does not).
 - Half as many flops as LU !
 - Only calculate L not U

Notes

Cholesky Factorization

Listing 2: Cholesky

```
1  given A
2  output L
3
4  for k = 1...n
5       $\ell_{kk} = \left( a_{kk} - \sum_{s=1}^{k-1} \ell_{ks}^2 \right)^{1/2}$  %SQUARE ROOT OF D ENDS UP ON DIAGONAL
6
7      for i = k + 1...n
8           $\ell_{ik} = \left( a_{ik} - \sum_{s=1}^{k-1} \ell_{is} \ell_{ks} \right) / \ell_{kk}$ 
9      end
10 end
```

Why SPD?

In general, SPD gives us

- non singular
 - If $x^T A x > 0$, for all nonzero x
 - Then $\lambda > 0$ for all eigenvalues λ
- No pivoting
 - From algorithm, can derive that
$$|l_{kj}| \leq \sqrt{a_{kk}}$$
 - Elements of \tilde{L} do not grow with respect to A
 - *For short proof see book*
- Cholesky faster than LU
 - No pivoting
 - Only calculate \tilde{L} , not U

Why SPD?

A matrix is Positive Definite (PD) if for all $x \neq 0$ the following holds

$$x^T A x > 0$$

- For SPD matrices, use the Cholesky factorization, $A = \tilde{L}\tilde{L}^T$
- Cholesky Factorization
 - Requires no pivoting
 - Requires one half as many flops as LU factorization, that is only calculate \tilde{L} not L and U .
 - Cholesky will be more than *twice* as fast as LU because no pivoting means no data movement
- Use MATLAB's built-in `cho1` function for routine work

Motivation Revisted

Multiple right hand sides

- Solve $Ax = b$ for k different b vectors
- Using LU factorization, the cost is $\mathcal{O}(n^3) + \mathcal{O}(kn^2)$
- Using Gaussian Elimination, the cost is $\mathcal{O}(kn^3)$

If A is symmetric

- Save 50% of the flops with LDL^T factorization
- Save 50% of the flops and many memory operations with Cholesky ($\tilde{L}\tilde{L}^T$) factorization

See `time_LU_vs_Cholesky.m`

The Backslash Operator

Given an $n \times n$ matrix A , and an $n \times 1$ vector b the `\` operator performs a sequence of tests on the A matrix. MATLAB attempts to solve the system with the method that gives the least roundoff and the fewest operations.


When A is an $n \times n$ matrix:

- 1 MATLAB examines A to see if it is a permutation of a triangular system
If so, the appropriate triangular solve is used.
- 2 MATLAB examines A to see if it *appears* to be symmetric and positive definite.
If so, MATLAB attempts a Cholesky factorization and two triangular solves.
- 3 If the Cholesky factorization fails, or if A does not appear to be symmetric,
MATLAB attempts an LU factorization and two triangular solves.

Summary: Complexity of Linear Solves

- $Ax = b$
- diagonal system: $\mathcal{O}(n)$
- tridiagonal system: $\mathcal{O}(n)$ (bigger constant)
- m -band system: $\mathcal{O}(m^2n)$
- upper or lower triangular system: $\mathcal{O}(n^2)$
- full system with GE: $\mathcal{O}(n^3)$
- scaled partial pivoting adds $\mathcal{O}(n^2)$
- full system with LU: $\mathcal{O}(n^3)$
- LU back solve: $\mathcal{O}(n^2)$
- m different right-hand sides (when using LU): $\mathcal{O}(n^3 + mn^2)$

Summary: Complexity



Approximate solutions

So far, we are seeking “exact” (up to floating point error) solutions x^* to

$$Ax = b.$$

What if we only need an approximation \hat{x} to x^* ?

We would like some \hat{x} so that $\|\hat{x} - x^*\| \leq \epsilon$, where ϵ is some tolerance.

The Residual

We can't actually evaluate

$$e = x^* - \hat{x}$$

But...

For $x = x^*$

$$b - Ax \equiv 0$$

For $x = \hat{x}$

$$b - Ax \neq 0$$

We call $\hat{r} = b - A\hat{x}$ the *residual* (Recall from Lecture 7). It is a different way to measure the “error”. In fact

$$\begin{aligned}\hat{r} &= b - A\hat{x} \\ &= Ax^* - A\hat{x} \\ &= A\hat{e}\end{aligned}$$

How big is the residual?

For a given approximation, \hat{x} to x , how “big” is the residual $\hat{r} = b - A\hat{x}$?

- $\|r\|$ gives a magnitude
- $\|r\|_1 = \sum_{j=1}^n |r_j|$
- $\|r\|_2 = \left(\sum_{j=1}^n r_j^2\right)^{1/2}$
- $\|r\|_\infty = \max_{1 \leq j \leq n} |r_j|$

Notes

Approximating x ...

Suppose we made a wild guess to the solution x of $Ax = b$:

$$x^{(0)} \approx x$$

How do I improve $x^{(0)}$?

Ideally:

$$x^{(1)} = x^{(0)} + e^{(0)}$$

but to obtain $e^{(0)}$, we must know x^* . Not a viable method.

Ideally (another way):

$$\begin{aligned} x^{(1)} &= x^{(0)} + e^{(0)} \\ &= x^{(0)} + (x^* - x^{(0)}) \\ &= x^{(0)} + (A^{-1}b - x^{(0)}) \\ &= x^{(0)} + A^{-1}(b - Ax^{(0)}) \\ &= x^{(0)} + A^{-1}r^{(0)} \end{aligned}$$

An iteration

Again, the method

$$x^{(1)} = x^{(0)} + A^{-1}r^{(0)}$$

is nonsense since A^{-1} is needed.

What if we approximate A^{-1} ? Suppose $Q^{-1} \approx A^{-1}$ and is cheap to evaluate, then

$$x^{(1)} = x^{(0)} + Q^{-1}r^{(0)}$$

is a good step.

continuing...

$$x^{(k)} = x^{(k-1)} + Q^{-1}r^{(k-1)} \tag{1}$$

What kind of Q^{-1} do I need?

Rewrite:

$$x^{(k)} = x^{(k-1)} + Q^{-1}(b - Ax^{(k-1)})$$

This becomes

$$\begin{aligned} Qx^{(k)} &= Qx^{(k-1)} + (b - Ax^{(k-1)}) \\ &= (Q - A)x^{(k-1)} + b \end{aligned}$$

Or

$$x^{(k)} = Q^{-1}(Q - A)x^{(k-1)} + Q^{-1}b \tag{2}$$

Often (1) is more practical for programming, while (2) is useful to do math. There are exceptions however.

Choices, choices....

We want Q^{-1} to be simple/cheap to apply.
Ideally it will also be a “good” approximate of A^{-1} .

Let $A = D - C_L - C_U$ where D is diagonal, C_L and C_U are *strictly* lower and upper triangular.

Choices, choices....

We want Q^{-1} to be simple/cheap to apply.
Ideally it will also be a “good” approximate of A^{-1} .

Let $A = D - C_L - C_U$ where D is diagonal, C_L and C_U are *strictly* lower and upper triangular.

What if $Q = D$? Super cheap to invert, maybe a bad approximation. . . What if $Q = D - C_L$? Still kind of cheap, probably a better approximation . . .

Two Popular Choices

Example

Jacobi iteration approximates A with $Q = \text{diag}(A)$.

```
1  $x = x^{(0)}$ 
2
3  $Q = D$ 
4
5 for  $k = 1$  to  $k_{\max}$ 
6    $r = b - Ax$ 
7   if  $\|r\| \leq \text{tol}$ , stop
8
9    $x = x + Q^{-1}r$ 
10 end
```

Two Popular Choices

Let $A = D - C_U - C_L$, where C_U and C_L are the negative of the strictly upper and lower triangular, respectively, of A .

Example

Gauss-Seidel iteration approximates A with $Q = \text{lowertri}(A)$.

```
1  $x = x^{(0)}$ 
2
3  $Q = D - C_L$ 
4
5 for  $k = 1$  to  $k_{max}$ 
6    $r = b - Ax$ 
7   if  $\|r\| \leq tol$ , stop
8
9    $x = x + Q^{-1}r$ 
10 end
```

Why D and $D - C_L$?

Look again at the iteration

$$x^{(k)} = x^{(k-1)} + Q^{-1}r^{(k-1)}$$

Looking at the error:

$$x^* - x^{(k)} = x^* - x^{(k-1)} - Q^{-1}r^{(k-1)}$$

Gives

$$e^{(k)} = e^{(k-1)} - Q^{-1}Ae^{(k-1)}$$

or

$$e^{(k)} = (I - Q^{-1}A)e^{(k-1)}$$

or

$$e^{(k)} = (I - Q^{-1}A)^k e^{(0)}$$

Why D and $D - C_L$?

We want

$$\mathbf{e}^{(k)} = (I - Q^{-1}A)^k \mathbf{e}^{(0)}$$

to converge.

Our iteration converges when $\|I - Q^{-1}A\| < 1$, because

$$\begin{aligned}\|\mathbf{e}^{(k)}\| &= \|(I - Q^{-1}A)^k \mathbf{e}^{(0)}\| \\ &\leq \|I - Q^{-1}A\|^k \|\mathbf{e}^{(0)}\|\end{aligned}$$

goes to 0, as $k \rightarrow \infty$.

Notes

Matrix Norms (Recall Lecture 7)

What is $\|I - Q^{-1}A\|$?

- $\|A\|_1 = \max_{1 \leq j \leq n} \sum_{i=1}^n |a_{ij}|$
- $\|A\|_\infty = \max_{1 \leq i \leq n} \sum_{j=1}^n |a_{ij}|$

Other norms (you may be familiar with):

- $\|A\|_2 = \sqrt{\rho(A^T A)}$
- $\rho(A) = \max_{1 \leq j \leq n} |\lambda_j|$
- $\|A\|_2 = \rho(A)$ for symmetric A

Again, why do Jacobi and Gauss-Seidel work?

Jacobi, Gauss-Seidel (sufficient) Convergence Theorem

If A is diagonally dominant, then the Jacobi and Gauss-Seidel methods converge for any initial guess $x^{(0)}$.

Definition: Diagonal Dominance

A matrix is *diagonally dominant* if

$$|a_{ii}| > \sum_{j=1, j \neq i}^n |a_{ij}|$$

for all i .

Notes

Smart Jacobi Algorithm

The Jacobi algorithm uses the matrix representation:

$$x^{(k)} = D^{-1}(C_L + C_U)x^{(k-1)} + D^{-1}b$$

Or componentwise,

$$x_i^{(k)} = - \sum_{j=1, j \neq i}^n \left(\frac{a_{ij}}{a_{ii}} \right) x_j^{(k-1)} + \frac{b_i}{a_{ii}}$$

So each sweep (from $k - 1$ to k) uses $\mathcal{O}(n)$ operations per vector element (potentially $\mathcal{O}(n^2)$ total).

If, for each row i , $a_{ij} \neq 0$ for *at most* m values, each sweep uses $\mathcal{O}(mn)$ operations.

Smart Gauss-Seidel Algorithm

The Gauss-Seidel algorithm uses the matrix representation:

$$x^{(k)} = (D - C_L)^{-1} C_U x^{(k-1)} + (D - C_L)^{-1} b$$

Component-wise:

$$x_i^{(k)} = - \sum_{j=1}^{i-1} \left(\frac{a_{ij}}{a_{ii}} \right) x_j^{(k)} - \sum_{j=i+1}^n \left(\frac{a_{ij}}{a_{ii}} \right) x_j^{(k-1)} + \frac{b_i}{a_{ii}}$$

So again each sweep (from $k - 1$ to k) uses $\mathcal{O}(n)$ operations per vector element.

If, for each row i , $a_{ij} \neq 0$ for *at most* m values, each sweep uses $\mathcal{O}(mn)$ operations.

The difference is that in the Jacobi method, updates are saved (and not used) in a new vector. With Gauss-Seidel, an update to an element $x_i^{(k)}$ is used immediately.

Example:

Lets do Jacobi and Gauss-Seidel:

Example

Consider

$$A = \begin{bmatrix} 2 & -1 & 0 \\ -1 & 3 & -1 \\ 0 & -1 & 2 \end{bmatrix}, \quad b = \begin{bmatrix} 1 \\ 8 \\ -5 \end{bmatrix}, \quad x^{(0)} = \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix}$$

See `iterative_methods.m`